



R Tutorial   R Interface   Data Input   Data Management   Statistics  
Advanced Statistics   Graphs   Advanced Graphs

< ADVANCED  
STATISTICS

Generalized Linear Models

Discriminant Function

Time Series

Factor Analysis

Correspondence Analysis

Multidimensional Scaling

**Cluster Analysis**

Tree-Based Models

Bootstrapping

Matrix Algebra

Become a data  
scientist with R  
on DataCamp.



**Start For Free**

## R IN ACTION



**R in Action** (2nd ed) significantly expands upon this material. Use promo code **ria38** for a 38% discount.

# Cluster Analysis

R has an [amazing variety](#) of functions for [cluster analysis](#). In this section, I will describe three of the many approaches: hierarchical agglomerative, partitioning, and model based. While there are no best solutions for the problem of determining the number of clusters to extract, several approaches are given below.

## Data Preparation

Prior to clustering data, you may want to remove or estimate missing data and rescale variables for comparability.

```
# Prepare Data
mydata <- na.omit(mydata) # listwise deletion of missing
mydata <- scale(mydata) # standardize variables
```

## Partitioning

**K-means** clustering is the most popular partitioning method. It requires the analyst to specify the number of clusters to extract. A plot of the within groups sum of squares by number of clusters extracted can help determine the appropriate number of clusters. The analyst looks for a bend in the plot similar to a scree test in factor analysis. See [Everitt & Hothorn \(pg. 251\)](#).

```
# Determine number of clusters
wss <- (nrow(mydata)-1)*sum(apply(mydata,2,var))
for (i in 2:15) wss[i] <- sum(kmeans(mydata,
  centers=i)$withinss)
plot(1:15, wss, type="b", xlab="Number of Clusters",
  ylab="Within groups sum of squares")
```

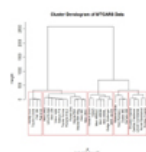
```
# K-Means Cluster Analysis
fit <- kmeans(mydata, 5) # 5 cluster solution
# get cluster means
aggregate(mydata,by=list(fit$cluster),FUN=mean)
# append cluster assignment
mydata <- data.frame(mydata, fit$cluster)
```

A robust version of **K-means** based on medoids can be invoked by using **pam()** instead of **kmeans()**. The function **pamk()** in the [fpc](#) package is a wrapper for **pam** that also prints the suggested number of clusters based on optimum average silhouette width.

## Hierarchical Agglomerative

There are a wide range of hierarchical clustering approaches. I have had good luck with Ward's method described below.

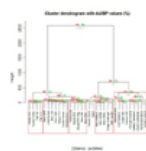
```
# Ward Hierarchical Clustering
d <- dist(mydata, method = "euclidean") # distance matrix
fit <- hclust(d, method="ward")
plot(fit) # display dendrogram
groups <- cutree(fit, k=5) # cut tree into 5 clusters
# draw dendrogram with red borders around the 5 clusters
rect.hclust(fit, k=5, border="red")
```



[click to view](#)

The `pvclust()` function in the [pvclust](#) package provides p-values for hierarchical clustering based on multiscale bootstrap resampling. Clusters that are highly supported by the data will have large p values. Interpretation details are provided [Suzuki](#). Be aware that [pvclust](#) clusters columns, not rows. Transpose your data before using.

```
# Ward Hierarchical Clustering with Bootstrapped p values
library(pvclust)
fit <- pvclust(mydata, method.hclust="ward",
  method.dist="euclidean")
plot(fit) # dendrogram with p values
# add rectangles around groups highly supported by the data
pvrect(fit, alpha=.95)
```



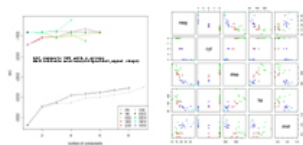
[click to view](#)

## Model Based

Model based approaches assume a variety of data models and apply maximum likelihood estimation and Bayes criteria to identify the most likely model and number of clusters. Specifically, the `Mclust()` function in the [mclust](#) package selects the optimal model according to BIC for EM initialized by hierarchical clustering for parameterized Gaussian mixture models. (pewh!). One chooses the model and number of clusters with the largest BIC. See [help\(mclustModelNames\)](#) to details on the model chosen as best.

```
# Model Based Clustering
library(mclust)
fit <- Mclust(mydata)
```

```
plot(fit) # plot results
summary(fit) # display the best model
```



[click to view](#)

## Plotting Cluster Solutions

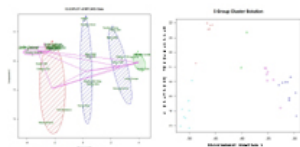
It is always a good idea to look at the cluster results.

```
# K-Means Clustering with 5 clusters
fit <- kmeans(mydata, 5)

# Cluster Plot against 1st 2 principal components

# vary parameters for most readable graph
library(cluster)
clusplot(mydata, fit$cluster, color=TRUE, shade=TRUE,
         labels=2, lines=0)

# Centroid Plot against 1st 2 discriminant functions
library(fpc)
plotcluster(mydata, fit$cluster)
```



[click to view](#)

## Validating cluster solutions

The function `cluster.stats()` in the [fpc](#) package provides a mechanism for comparing the similarity of two cluster solutions using a variety of validation criteria (Hubert's gamma coefficient, the Dunn index and the corrected rand index)

```
# comparing 2 cluster solutions
library(fpc)
cluster.stats(d, fit1$cluster, fit2$cluster)
```

where `d` is a distance matrix among objects, and `fit1$cluster` and `fit2$cluster` are integer vectors containing classification results from two different clusterings of the same data.

## To Practice

Try the clustering exercise [in this introduction to machine learning course](#).

Copyright © 2017 Robert I. Kabacoff, Ph.D. | Sitemap